

INTRODUCTION TO DOCKER AND K8S

Michael Bright, Kjell van Dam & Pascal Van Dam

April 28, 2021



INTRODUCTION

TRAINER & STUDENT INTRODUCTION

- » Introduce yourself shortly
- » Do you have any experience with
 - Containerization
 - Docker
 - Kubernetes
 - Linux
- » Certification ?

This course:

- » Is developed for Linux professionals that would like to know more about running and managing workloads in Docker & Kubernetes on Linux and in the cloud.
- » Will introduce you to the core concepts of Docker and Kubernetes
- » Enables you to manage the Kubernetes system on a high level
- » Enables you to develop simple containerized workloads

CKA

Certified Kubernetes Administrator

CKAD

Certified Kubernetes Developer

CKS

Certified Kubernetes Security Specialist

*See <https://training.linuxfoundation.org>
for more information and the “Certification preparation guide”*

- » 2 hours, lab-based exam
- » All about managing K8S clusters

- » 2 hours, lab-based exam
- » Focusses on developing containerized workload and transforming workload

- » 2 hours, lab-based exam
- » Focusses on security aspects of K8S and running containerized workload

- » Introduction
- » Why containers?
- » Container Components
- » Docker introduction
- » Installing Docker
- » Building Container Images
- » Docker Registries

- » Docker Volumes
- » Docker Configs & Secrets
- » Docker Compose & Stack
- » Docker Swarm
- » Introduction to Kubernetes
- » Kubernetes architecture
- » Installing K8S using kubeadm
- » First steps on K8S

- » Kubernetes PODs
- » Kubernetes Deployments
- » Kubernetes DaemonSets
- » Kubernetes Jobs
- » Kubernetes CronJobs
- » Volumes
- » ConfigMaps and Secrets

GETTING STARTED WITH CONTAINERS

GETTING STARTED WITH CONTAINERS

WHY CONTAINERS

- » In the beginning each application had its own server
New application needed? → New server deployment
- » Advantages:
 - Ultimate isolation of applications
 - Very secure
 - Easy to tune the OS for one single application
- » Disadvantages:
 - Very expensive
 - Very inefficient (low utilization)
 - Not agile; long time to market

AND THEN THERE WAS VIRTUALIZATION

- » Hypervisor technology introduced the possibility to run multiple operating systems on one server
- » Advantages:
 - Much better server utilization (>80%)
 - Faster time to market
 - More agile
 - Well isolated (security & manageability)
- » Disadvantages:
 - Still high CAPEX and OPEX costs for OSES
 - Configuration management challenge (VM Sprawl)
 - Still not fast and agile enough
 - Still a limited number of applications on one OS/server

- » Containers allow multiple applications to run in a standardized isolated environment within one single OS
- » Advantages:
 - Best utilization/density
 - Less overhead
 - Very agile
 - Blazingly fast time to market
- » Disadvantages:
 - Less isolation compared to virtual machines
 - Access to physical/virtual hardware is difficult

WHY DO I NEED CONTAINERS?

Containers are the answer to business desire for having

- » Better hardware utilization
- » Increasingly faster times-to-market for the applications
- » Reduction of risk and complexity while deploying
- » A uniform industry standard way of deploying applications
- » Having more control on the application environments
- » The desire to be as agile as possible

WHAT ARE CONTAINERS



- » Resource partition technology
- » Very light weight
- » An Industry Standard
- » Revolutionizing working with applications:
 - Agile workflow from development to production
 - Integration with version control systems
 - Independent features
 - Automatic testing
 - Rapid failback

You may now start with the following labs:

- » 1.1 Docker installation
- » 1.2 Running Containers

GETTING STARTED WITH CONTAINERS

CONTAINER COMPONENTS

- » Namespaces
- » Cgroups
- » Storage
- » Networking
- » Security Framework

- » Partition and isolate a global resource
- » Processes in the name space see their own isolated instance
- » Similar to `chroot`, extended to other global resources
- » Heavily used by containers

- » Mount namespace view
- » Chroot is used as the foundation
- » Each container has it's own root filesystem (/)

- » PID namespace view - used in each container
- » Each container has it's own PID 1
- » Outside of the container this will be a different PID
- » None of the containers can see, start or kill processes on other containers or on the container host

- » User name space keeps a dedicated isolated user database
- » Each container has its own user database. For instance UID 0 (root) in the container is not UID 0 outside of the container

- » **Inter Process Communications (IPC):** Partitioning and Isolation of SysV IPC like shared memory, semaphores and message queues
- » **Networking:** Partitioning and Isolation of the networking stack Processes within the namespace have the experience of seeing their own network stack, independent of the container host stack
- » **UTS:** (Unix Timesharing System) Allow processes in a namespace to have their own hostname and (NIS) domain name



- » **Time namespace:** Gives namespaced processes their own view on a (monotonic) clock

- » By default on Linux and UNIX all processes are equal But some processes are more equal than others
- » CGroups are resource pools to share and limit resources
- » Processes are wrapped in CGroups or Child CGroups
- » CGroups available for cpu, blkio, mem, network and devices
- » Heavily used in systemd and in container technology

- » Linux users are privileged user (root) or non-privileged
 - No restrictions apply to root
 - Regular users have restricted possibilities
- » To get enough permissions, a process is often started as root
- » Capabilities: allow fine-grained elevated privileges to non-privileged users

- » Containers utilize capabilities to get access to privileged functions
- » Some examples:
 - `CAP_NET_SERVICE_BIND`: to allow binding of network ports < 1023
 - `CAP_MKNOD` to allow creation of device nodes
 - `CAP_CHOWN` to allow changing ownership of files
- » See `man 7 capabilities` for more information

Containers need storage to:

- » Store the container images (filesystem, binaries, libs)
- » Store the persistent data (*optional*)

STORING CONTAINER IMAGES (1)

- » Container images are stored in layers
- » Copy-on-write (CoW) mechanism
- » Each layer stacks upon the previous layer
- » Only the top layer is writable
- » Different vendors are using different storage drivers:
 - Devicemapper (direct-lvm) → RHEL, Fedora and CentOS
 - AUFS → Ubuntu
 - BTRFS → SLES, OpenSUSE LEAP

- » CoW makes starting and restarting containers very fast
- » Changes are light-weight in the container images (stacked)
- » Designed for storage efficiency, not speed

You may now start with the following labs:

- » 1.3 Create your own images
- » 1.4 Experiments with persistency


You may now start with the following labs:

- » 1.3 Create your own images
- » 1.4 Experiments with persistency

GETTING STARTED WITH CONTAINERS

DOCKER INTRODUCTION


- » Initial release in 2013
- » Open-Source but strictly managed by Docker Inc
- » On 13-11-2019 split up into Docker and Mirantis
- » Docker: focus on developer workflow with Docker
- » Mirantis: focus on Docker Enterprise (MKE)



- » CLI with GIT like interface
- » Container Images based on lightweight layers
- » Docker Hub, Docker Registry, Docker Datacenter
- » Orchestration and clustering possible with Swarm or 3th party software



- » Docker CLI
- » Docker Engine
- » Containerd
- » RunC
- » Containers



- » **Container image:** The contents / package that can be run in a container. base image + your application.
- » **Base image:** Only the fresh OS, not additional layers
- » **Image layer:** Each change will result in a new layer stacked upon the container image
- » **Container:** Running instance of an image; e.g. your running application / service
- » **Container Host or Docker Host:** The host where Docker containers are running
- » **Docker Hub:** Generic image store on the Internet
- » **Docker Registry:** An (optional) on-premise image store

- » Docker can be installed from the distribution repositories or from the Docker website
- » Two Editions:
 - Docker CE (Community Edition)
 - Docker EE (Enterprise Edition)
 - Distro specific version: RHEL7 -> Docker 1.13.x
- » Pre-requisites:
 - Min. 2Gb of RAM
 - Min. 3Gb of storage for container images
 - Optional: storage for persistent volumes

RHEL based distributions:

- » Install dependencies:
 - device-mapper-persistent-data
 - lvm2
 - RHEL/Centos 7.5 and further -> Overlay2FS
- » Add the repository using `yum-config-manager`
- » Install docker-ce with `yum`

Debian based distributions:

- » Install `apt-transport-https` to be able to access https repositories
- » Import the the GPG key, using `apt-key`
- » Add the repository using `add-apt-repository`
- » Install `docker-ce` with `apt`

1. Start and enable Docker:

```
systemctl enable --now docker
```

2. Add the user(s) that need to administer Docker to the docker group:

```
sudo usermod -aG docker <user>
```

3. Logout and login again
4. Any user belonging to the docker group can run the Docker commands
5. Verify Docker status with such a user:

```
docker info
```

- » Pull an image from Docker Hub into a container and run it:

```
docker run <image name>
```

- » Only pulling an image from Docker Hub

```
docker pull <image name>
```

- » Start a container

```
docker start <container name>
```

- » Stopping a container

```
docker stop <container name>
```

- » Removing a container

```
docker rm [-f] <container name>
```

- » Restart a container

```
docker restart <container name>
```


- » Search images by name on Docker Hub:

```
docker search <name>
```

- » Filter on official images:

```
docker search --filter "is-official=true"
```

- » Other filters:

- Rating: `stars=`
- Automated builds `is-automated=true|false`

- » Docker images can be pulled from Docker hub or created by hand:

```
docker build
```

- » This command build images, using a Dockerfile as an input file and built the images given the commands in this file
- » The Dockerfile contains the commands how to build the image

```
FROM ubuntu:latest
```

```
MAINTAINER Pascal van Dam (pascal@yunix.org)
```

```
RUN apt-get update
```

```
RUN apt-get install -y python python-pip wget
```

```
RUN pip install Flask
```

```
WORKDIR /home
```

```
ADD hello.py /home/hello.py
```

```
CMD ["python", "./hello.py"]
```

- » Build the image:

```
docker build . --tag dockertest:latest
```

- » Verify if the built image is locally available now:

```
docker image ls | grep dockertest
```

- » Try it:

```
docker run dockertest:latest
```

- » Next steps:

- Push it to a local registry or Docker Hub

- » Docker images will be built layer by layer
- » Each command will generate a layer
- » Tip: Limit the number of layers by grouping commands with compound statements

Dockerfile

```
FROM node
```

```
WORKDIR /app
```

```
COPY package.json .
```

```
RUN npm install express -save && npm install && mkdir /app/public
```

```
COPY helloworld.js /app/
```





```
COPY public/* /app/public/
```

```
EXPOSE 8081
```

```
CMD [ "node", "helloworld.js" ]
```

- » At container startup you can run the container in interactive mode:

```
docker run -it <image> --name <container>
```

- » Leave the interactive mode using  + ,  + 
- » Invoke a Bash shell, when the container is already running:

```
docker exec -it <container name> /bin/bash
```

- » Get the stdout/stderr info from the containers console:

```
docker logs <container>
```

- » Deep dive into the container configuration, to get information about:

- Network information
- Volume information
- Image information

```
docker inspect <container>
```


GETTING STARTED WITH CONTAINERS

DOCKER REGISTRIES

- » It is possible to store and retrieve images from a registry
- » Docker Hub `https://hub.docker.com`
- » A third Party registry (ACR, ECR etc)
- » a private registry (docker, harbor)

- » Docker Hub is a registry owned by Docker INC.
- » Can be used for publicly and privately hosted container images
- » Information about containers can be found on Docker Hub website:
 - Dockerfile
 - How to configure and the use image
 - Tips and tricks
 - Related images

- » Search for an image:

```
docker search nginx
```

- » More detailed search:

```
docker search --filter "is-official=true" --no-trunc nginx
```

- » Download and run:

```
docker run --name nginxc01 -p 80:80 nginx
```

» Create a free account on <https://hub.docker.com>

» Storing an image on Docker Hub:

1. Log in with your Docker Hub account:

```
docker login -u <username> -p <password>
```

2. Tag your image for storing on Docker Hub:

```
docker tag <id> <accountname>/<image>:versiontag
```

3. Push the image

```
docker push
```

4. Verify that both original and tagged images are listed:

```
docker images
```

- » Secure or insecure
- » Authentication is an option for secure registries
- » Protocol used is HTTPS
- » There are alternatives for the Docker Registry like Harbor or Quay.io

What is needed for a simple insecure registry?

- » A docker node to run the registry container on
- » A TCP port on which the registry will listen
- » Persistent storage to store the container images
- » The **registry:2** image from Docker Hub

CREATE A PRIVATE REGISTRY

```
docker run --detach \  
  --restart=always \  
  --name registry \  
  --publish 5000:5000 \  
  --volume /srv/registry:/var/lib/registry \  
  registry:2
```


To use an insecure registry we have to declare it as 'trusted' in the `/etc/docker/daemon.json` file. After that the docker daemon needs to be restarted.

`/etc/docker/daemon.json`

```
{  
  "insecure-registries" : [ "st99node01:5000", "st99node01.itgildelab.net:5000"]  
}
```

restart docker daemon

```
sudo systemctl restart docker
```

What is needed for a secure registry?

- » A Docker node to run the registry container on
- » A TCP port on which the registry will listen
- » Persistent storage to store the container images
- » SSL certificate(s) and key
- » Certificate must be added to `/etc/docker/certs.d`

First create the needed certificate and private key

```
mkdir certs
cd certs

openssl req -new -sha256 -newkey rsa:4096 -x509 -sha256 \
  -nodes -days 365 -out registry.crt -keyout registry.key \
  -subj "/C=NL/ST=LB/O=Acme, Inc./CN=registry.itgilde.lab"
```

CREATE A SECURE REGISTRY

Spin-up the container using the created certificate and key.

```
docker run -d \  
  --restart=always \  
  --name registry \  
  -v ${PWD}/certs:/certs \  
  -e REGISTRY_HTTP_ADDR=0.0.0.0:443 \  
  -e REGISTRY_HTTP_TLS_CERTIFICATE=/certs/registry.crt \  
  -e REGISTRY_HTTP_TLS_KEY=/certs/registry.key \  
  -p 443:443 \  
  -v /srv/registry:/var/lib/registry \  
  registry:2
```

On every docker client, create a directory under `/etc/docker/certs.d` and place the certificate in it. If the port is unequal to 443 please also specify the port in the URL directory name.

```
mkdir -p /etc/docker/certs.d/st99node01.itgildelab.net  
cp certs/st99node01.itgildelab.net.crt /etc/docker/certs.d/st99node01.itgildelab.net
```

You may now start with the following labs:

» 1.5 Creating registries

GETTING STARTED WITH CONTAINERS

VOLUMES AND MOUNTS

Since the early days of Docker there has been the concept of bind mounts.

- » A file or directory from the host filesystem is mounted in the container
- » Has limited functionality compared to volumes
- » Use `-v` or `--volume`

EXAMPLE: BIND MOUNT WITH VOLUME OPTION

```
docker run --detach --interactive --tty \  
  --name devtest \  
  --volume $(pwd)/html:/usr/share/nginx/html \  
  nginx:latest
```

- » The `--volume` is only supported in stand-alone containers
- » `--mount` works for stand-alone containers and in swarm mode
- » In general `--mount` is more explicit and verbose

EXAMPLE: BIND MOUNT WITH MOUNT OPTION

```
docker run --detach --interactive --tty \  
  --name devtest \  
  --mount type=bind,source=$(pwd)/html, \  
    target=/usr/share/nginx/html \  
  nginx:latest
```

Volumes are the preferred way to supply persisting storage to containers.

- » Volumes are easier to backup than bind mounts
- » Volumes can be managed using the Docker CLI
- » Volumes work on Linux and Windows containers
- » Volumes can be shared in a safer way between containers
- » Volume drivers are available for external storage provisioning
- » New volumes can be pre-populated by a container

- » Create a volume:

```
docker volume create <volume name>
```

- » List volumes:

```
docker volume ls
```

- » More details of a volume:

```
docker volume inspect <volume name>
```

- » Remove a volume:

```
docker volume rm <volume name>
```

STARTING A CONTAINER WITH A VOLUME

If you start a container with a volume that does not exist yet, Docker will create it for you.

```
docker run -d \  
  --name devtest \  
  --mount source=myvol2,\  
    target=/usr/share/nginx/html \  
  nginx:latest
```

To create a shared NFS volume:

```
docker volume create \  
  --driver local --opt type=nfs \  
  --opt o=addr=st00node01,rw \  
  --opt device=:/mnt/pvs/st00pvol01/ nfsvol
```

```
docker run --rm -it \  
-v nfsvol:/data alpine
```


GETTING STARTED WITH CONTAINERS

DOCKER COMPOSE

- » Managed multiple related containers
- » Creates network connections between containers
- » Faciliates volume creation
- » Non-swarm: docker-compose
- » Swarm: docker stack

```
sudo curl -L \
  "https://github.com/docker/compose/releases/download/1.23.0"
-o /usr/local/bin/docker-compose
```

```
version: "3.9"

services:
  db:
    image: mysql:5.7
    volumes:
      - db_data:/var/lib/mysql
    restart: always
    environment:
      MYSQL_ROOT_PASSWORD: somewordpress
      MYSQL_DATABASE: wordpress
      MYSQL_USER: wordpress
      MYSQL_PASSWORD: wordpress
```

GETTING STARTED WITH CONTAINERS

DOCKER SWARM

- » Docker Swarm "classic" - early version
- » Swarmkit - toolkit for building HA apps
- » Docker "Swarm Mode" - Docker native orchestration

- » Docker native orchestrator
- » No further developments by Mirantis
- » Advised orchestrator is: K8S

- » Docker swarm runs services not containers
- » Deployment patterns
 - Replicas
 - Global replicas
- » Advised orchestrator is: K8S


```
version: "3.9"
services:
  db:
    image: mysql:5.7
    volumes:
      - db_data:/var/lib/mysql
    restart: always
    environment:
      MYSQL_ROOT_PASSWORD: somewordpress
      MYSQL_DATABASE: wordpress
      MYSQL_USER: wordpress
      MYSQL_PASSWORD: wordpress

  wordpress:
```

GETTING STARTED WITH CONTAINERS

CONFIGS

- » Separates code and config in the container
- » Config is added at runtime to the container
- » Configs are for non-sensitive information
- » BASE64 encode, not encrypted
- » Configs will be mounted inside the container
- » Only works in SWARM mode

```
docker config create config-v1 index.html  
docker config create config-v2 alternate-index.html
```

```
docker config inspect config-v2
```

```
docker service create --name web -p 80:80 \  
  --config source=config-v1,target=/usr/share/nginx/html/index.html \  
  nginx:latest
```

```
docker service update web \  
  --config-rm config-v1 \  
  --config-add source=config-v2,target=/usr/share/nginx/html/index.html \  
  --update-order start-first
```

GETTING STARTED WITH CONTAINERS

SECRETS

- » Like Configs, separate code and config in the container
- » Secrets are added at runtime to the container
- » Secrets are for non-sensitive information
- » Secrets are SHA256 encrypted (@rest and in transit)
- » Secrets will be mounted inside the container
- » Only works in SWARM mode

```
docker secret create secret-v1 index.html  
docker secret create secret-v2 alternate-index.html
```

```
docker secret inspect secret-v2
```

```
docker service create --name web -p 80:80 \  
  --secret source=secret-v1,target=/usr/share/nginx/html/index.html \  
  nginx:latest
```

```
docker service update web \  
  --secret-rm secret-v1 \  
  --secret-add source=secret-v2,target=/usr/share/nginx/html/index.html \  
  --update-order start-first
```

GETTING STARTED WITH CONTAINERS

CLEANING UP THE DOCKER ROOM

Clean up dangling images

```
docker image prune
```

Clean up all unused images

```
docker image prune -a
```

Prune images which are older dan 1d

```
docker image prune -a --filter "until=24h"
```

Removing the container after exit

```
docker run --rm -detach --name <name> <image>
```

Clean up old containers

```
docker container prune
```


Clean up unreferenced volumes

```
docker volume prune
```

Label a volume

```
docker volume create --label <label> <volume name>
```

Clean up volumes that do not have a specific label

```
docker volume prune --filter "label!=<label>"
```

```
# Clean up unreferenced networks
```

```
docker network prune
```

CLEANING UP ALL DOCKER OBJECTS

```
# Clean up all unreferenced docker objects except volumes  
docker system prune
```

```
# Clean up all unreferenced docker objects including volumes  
docker system prune --volumes
```

CONTAINERS - NEW TOOLS ON THE BLOCK

CONTAINERS - NEW TOOLS ON THE BLOCK

NEW TOOLS ON THE BLOCK

Tools to replace docker tooling



podman



buildah



skopeo



- » Replacement for docker client
- » Almost a drop-in replacement
 - Accepts all docker commands and options
 - Extra options to remove all containers
 - Delegates `docker build` to buildah
- » OCI compliant



- » `podman pull pamvdam/pyco2:v4`
- » `podman build . --tag mycontainer:alpha1`
- » `podman rmi --all`
- » `podman rm --all --force`
- » This one is different:
 - `podman push`



buildah

- » Replaces docker build
- » Dockerfile compatability mode (bud)
- » No caching
- » Single commit at the end
- » No docker daemon used
- » Buildah native mode using shell scripts



skopeo

- » Tool for registry operations
- » Copying of images between repos
- » Transforming container formats

INTRODUCTION TO KUBERNETES

INTRODUCTION TO KUBERNETES

HISTORY

- » Started as GOOGLE project Borg.
- » Opensourced and released as Kubernetes
- » in Ancient greek: κυβερνήτης
 - Meaning: Helmsman, navigator, pilot
- » Google Project Seven
- » Founded by Joe Beda, Brendan Burns and Craig McLuckie
- » Maintained by the Cloud Native Computing Foundation (CNCF)
- » Popular referenced as K8S. ('Kates')

Incorporated in many solutions:

- » RedHat OpenShift
- » Rancher 2
- » MESOSPHERE - DC/OS
- » Azure: Azure Kubernetes Service
- » EKS: Elastic Kubernetes Service
- » GKE: Google Kubernetes Engine
- » IKS: IBM Kubernetes Service

INTRODUCTION TO KUBERNETES

KUBERNETES ARCHITECTURE

- » Master/slave architecture
- » Kubernetes Control Plane v.s. Worker Nodes
- » Composed of Building Blocks (Primitives)
 - Deploying applications
 - Maintaining applications
 - Scaling applications
- » Loosely coupled
- » All revolving around the Rest API
- » Extensible by API, containers and extensions.

- » Pods
- » Labels and selectors
- » Controllers
- » Services

- » Basic scheduling unit in Kubernetes
- » Contains one or more containers that are scheduled together
 - Guaranteed to be co-located on the same host
 - Can share resources together
- » Has a unique IP address
- » Share network stack and volumes
- » Can be managed manually using the rest API or by a controller

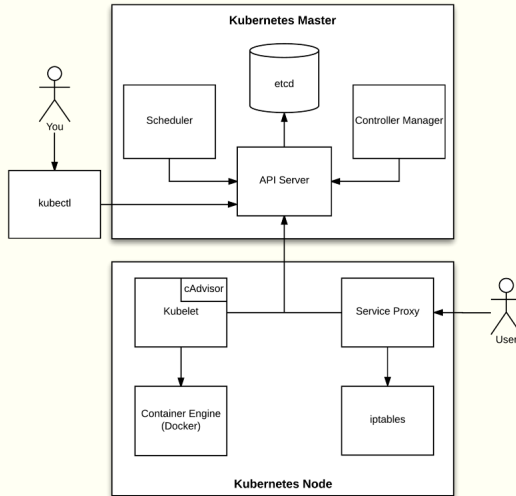
- » Labels are key-value pairs that can be attached to API objects like
 - Pods
 - Nodes
- » Label selectors are queries against labels
- » Example use cases:
 - Select to which pods traffics is routed to.
 - Select which pods get updated/scaled up/down etc.
- » Always use labels!

- » Managed by the Control Manager
- » A control loop that watches the shared state
- » Makes changes to move the current state towards the desired state
- » Example controllers:
 - ReplicaSet controller
 - DaemonSet controller
 - Job Controller

- » Logical set of PODs
- » Provides a single IP address and DNS name by which PODs can be accessed
- » Helps with LoadBalancing
- » Types of services
 - ClusterIP: access only from within the cluster
 - NodePort: access from outside the cluster on a static port
 - LoadBalancer: Uses cloud providers' Load Balancer facility


- » API server
- » ETCD
- » Scheduler
- » Controller Manager

KUBERNETES CONTROL PLANE - MASTER

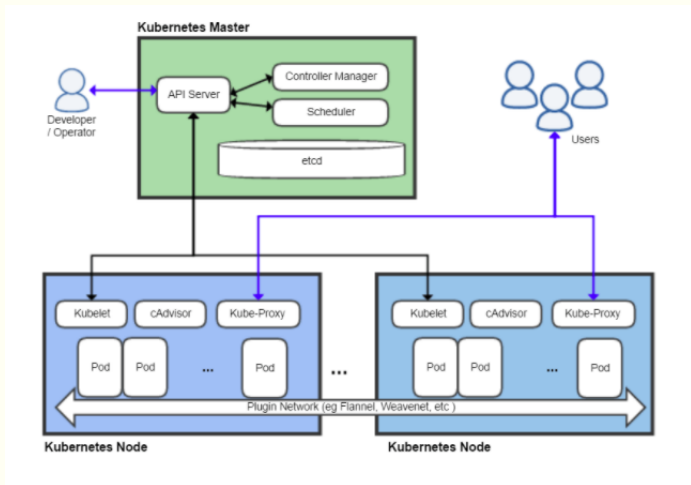


- » Kubelet - Controls state/manages containers
- » Container - contains the application
- » Kube-proxy - routes IP traffic to container


KUBERNETES WORKER NODE - NODE




KUBERNETES ARCHITECTURE



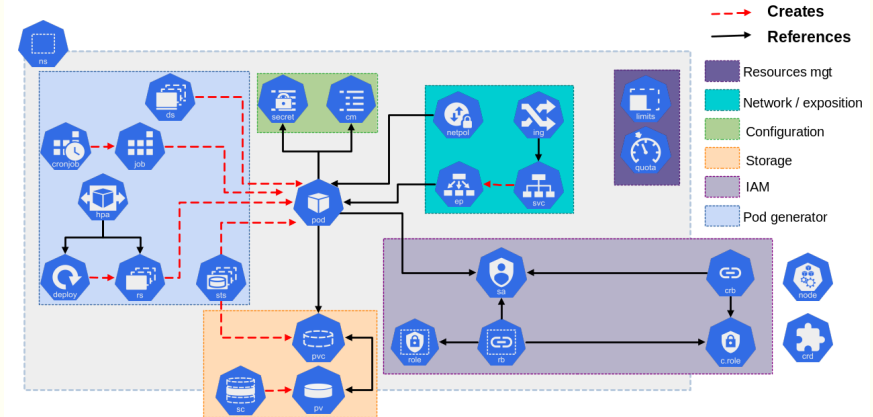
KUBERNETES ARCHITECTURE



EXAMPLE OF POD CREATION WORKFLOW



MAP OF MAIN K8S RESOURCES



- » PO - POd
- » RS - ReplicaSet
- » DEPLOY - DEPLOYment
- » HPA - Horizontal Pod Autoscaler
- » STS - StaTeful Set
- » CRJ - CronJob
- » JO - Job
- » SVC - SerViCe
- » CRD - Custom Resource Definition
- » RC - Replica Controller (deprecated)

- » EP - End Point
- » PV - Persistent Volume
- » PVC - Persistent Volume Claim
- » SC - Storage Class
- » CM - ConfigMap
- » SECRET - SECRET
- » DS - Daemon Set
- » NETPOL - NEtwork Policy

INTRODUCTION TO KUBERNETES

INSTALLING KUBERNETES

- » Fully from scratch
- » Minikube
- » Kubeadm
- » Using a Kubernetes service on a public cloud

- » Works on so called bare metal servers
- » Supports the latest kubernetes version
- » Can create single and multi master K8S clusters
- » Can facilitate upgrade to newer K8S version

- » Docker engine (Docker CE 19.03 recommended) installed and running
- » Swap must be turned off
- » At least 2GiB of RAM

- » Disable swap (*Don't forget to edit /etc/fstab*)

```
sudo swapoff -a
```

- » Install the Docker container runtime and start the engine as discussed in the Docker introduction
- » Configure systemd as the recommended driver for Docker

SYSTEMD AS DOCKER DRIVER (1)

```
su -
```

```
cat > /etc/docker/daemon.json <<EOF
```

```
{  
  "exec-opts": ["native.cgroupdriver=systemd"],  
  "log-driver": "json-file",  
  "log-opts": {  
    "max-size": "100m"  
  },  
  "storage-driver": "overlay2"  
}  
EOF
```

Directory for control files

```
mkdir -p /etc/systemd/system/docker.service.d
```

Restart docker.

```
systemctl daemon-reload
```

```
systemctl restart docker
```

- » Import the GPG key:

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg \  
| sudo apt-key add -
```

- » Add the repository (Debian and Ubuntu, all versions)

```
sudo apt-add-repository "deb http://apt.kubernetes.io/ \  
kubernetes-xenial main"
```

- » Install the software:

```
sudo apt update  
sudo apt install -y kubeadm kubect1 kubelet
```

Configure the master and define the subnet:

```
kubeadm init --pod-network-cidr <private subnet>
```

Please save the output of this command!

In case of reported issues by kubeadm

- » If `kubeadm` report sizing issues, add the parameter:
`--ignore-preflight-errors=<list of errors>`
- » If you are using another Docker registry, add the parameter:
`--image-repository <url>`

As a normal user (non-root) execute:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You may now start with Lab 2.1 in your labguide.

Check the current state:

```
kubectl get nodes
```

The master node will be listed a 'Not Ready', because the network is not configured.

- » Kubernetes is an orchestrator for containers
- » It doesn't manage networks
- » You'll need a Container Network Interface (CNI)
- » Many CNI's are available, such as:
 - Calico
 - Weave
 - Flannel

On Azure Calico will not work properly so we will use CANAL

```
kubectl apply -f \
  https://docs.projectcalico.org/v3.10/manifests/calico.yaml
```

On Azure Calico will not work properly so we will use CANAL

```
kubectl apply -f \  
    https://docs.projectcalico.org/v3.10/manifests/canal.yaml
```

- » Confirm that all of the pods are running:

```
watch kubectl get pods --all-namespaces
```

- » And review the master availability:

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
debian	Ready	master	47h	v1.19.3

- » Execute the same procedure as for the Master
- » Join the Kubernetes Master, using the saved output from the Master installation.

```
kubeadm join --token <token> <master-ip>:<master-port> \  
--discovery-token-ca-cert-hash sha256:<hash>
```

VERIFY THE CLUSTER STATUS

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
st00node01	Ready	master	47h	v1.19.3
st00node02	Ready	<none>	47h	v1.19.3
st00node03	Ready	<none>	46h	v1.19.3

MANAGING K8S

MANAGING K8S

FIRST STEPS ON K8S

- » In K8S we don't run containers, we run PODs
- » Use `kubect1 run` to start a POD.
- » This is the adhoc use of Kubernetes
- » This functionality will be deprecated.

K8S 1.19 and later - runs a standalone Pod

```
kubectl run nginx --image nginx:latest --port 80
```

K8S 1.18 and earlier - creates a single Pod Deployment

```
kubectl run nginx --image nginx:latest --port 80 --replicas=1
```

Use `kubectl get pods` to examine the results

```
kubectl get pods
```

```
kubectl get pods --namespace default
```

```
kubectl get pods --all-namespaces
```

```
kubectl get pods --namespace kube-system
```

MORE WAYS TO EXAMINE PODS

```
kubectl get pods -o wide
```

```
kubectl describe pod
```



```
kubectl delete pod <pod id>
```

```
kubectl get rs nginx  
kubectl describe rs nginx
```

```
kubectl scale --replicas=3 rs nginx  
kubectl get pods -o wide  
kubectl delete pod <pod-id>  
kubectl get events | head -10
```

```
kubectl delete rs nginx  
kubectl get pods -o wide
```

MANAGING K8S

INTRODUCTION TO KUBECTL

- » Kubectl is your one-stop-shop tool for managing K8S clusters
- » User interface of `kubectl` is very intuitive
- » General syntax: `kubectl <verb> <object> <options>`

```
kubectl get pods -o wide
```

```
kubectl get pods -o yaml
```

```
kubectl describe deployment
```

Kubectl knows the following verbs for read-like actions

- » describe

- » get

```
kubectl describe deployment mydeployment
```

```
kubectl get nodes
```

```
kubectl get pod -o wide
```

Kubectl knows the following verbs for create-like actions

» create

» run

```
kubectl create -f mydeployment.yml
```

```
kubectl run nginx --image=nginx --port=80
```

```
kubectl create deployment web --image http:latest --port 80 --replicas 1
```


Kubectl knows the following verbs for deleting objects

» delete

```
kubectl delete pods -l myapp  
kubectl delete svc nginx-service  
kubectl delete ns test
```

Kubectl knows the following verbs for updating objects

- » set
- » label
- » annotate
- » scale
- » edit

```
kubectl scale svc nginx-serice --replicas=3
kubectl edit deployment myapp
kubectl set image deployment.v1.apps/nginx-deployment
    nginx=nginx:1.9.1
```

Kubectl knows the following objects

- » nodes
- » pods
- » deployments
- » services
- » rs to manage replica sets
- » And a whole lot more...

```
kubectl expose deployment q10rv2 --type NodePort --port 5000
```

- » Managing K8S using imperative commands.
- » Managing K8S using imperative object configuration.
- » Managing K8S using declarative object configuration.

- » Easiest way to operate your cluster, good for starting
- » With `kubectl` you can operate directly on live objects
- » Useful for one-off tasks

```
kubectl run nginx --image nginx:1.7.1 --port=80  
kubectl set image deployment nginx nginx=nginx:1.9.1  
kubectl scale deployment nginx --replicas=1
```

IMPERATIVE CONFIGURATION (1)

- » More difficult, but more powerful
- » A YAML file that describes the new object or how it should be altered
- » Describe the desired state of the object(s) and have the controllers sort it out
- » Create the object(s):

```
kubectl create -f <yaml cfg file>
```
- » Delete the object(s):

```
kubectl delete
```
- » Create / Update the object(s)

```
kubectl apply
```

```
kubectl get deployment nginx -o yaml > nginx.yml  
kubectl delete -f nginx.yml  
kubectl create -f nginx.yml  
kubectl replace -f nginx.yml  
kubectl create -f http://myrepo.itgilde.lab/calico.yml
```

- » Configuration to reach the desired state
- » Complex to manage and to design, but very powerfull
- » Describe the desired state using a directory of manifest files and have K8S controllers figure it out
- » Will be extensively discussed in the Advanced Kubernetes Course

```
kubectl apply -f config/  
kubectl apply -R -f config/  
kubectl diff -R -f config/
```


- » Get help

```
kubectl get pods --help
```

- » Show logs of a pod

```
kubectl logs <pod>
```

```
kubectl logs <pod> -c <container>
```

- » Explain the yaml config file format for an object

```
kubectl explain <object>
```

```
kubectl explain pod  
kubectl explain pod.spec  
kubectl explain pod.spec.volumes
```

Together with `explain`, `kubectl run` can serve as an easy template generator.

» For a deployment:

```
kubectl run nginx --image nginx:1.15.1 --port 80
```

» For a bare pod:

```
kubectl run nginx --image=nginx --port 80 --restart=Never
```

» For a cron job:

```
kubectl run busybox --image=busybox \  
--schedule="* * * * *" --restart=OnFailure
```

Together with `explain`, `kubectl run` can serve as an easy template generator.

» For a deployment:

```
kubectl create deployment nginx --image nginx:1.15.1 --port
```

» For a bare pod:

```
kubectl run nginx --image=nginx --port 80
```

» For a job:

```
kubectl create job hello --image=busybox -- echo "Hello World"
```

OBJECT MANIFESTS

Each object in K8S can be described using a manifest

- » The manifest file is written in YAML
- » The manifest file has at least 4 parts: AKMS
 - A API (version)
 - K KIND (kind of object)
 - M METADATA (labels, annotations etc)
 - S SPEC (object specifications and attr)

EXAMPLE DEPLOYMENT MANIFEST (1)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: example-deploy
  annotations:
    deployment.kubernetes.io/revision: "1"
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: myhelloworld
```

EXAMPLE DEPLOYMENT MANIFEST (2)

```
spec:
  containers:
  - image: pamvdam/myhelloworld:v0.4
    imagePullPolicy: Always
    name: myhelloworld
    ports:
    - containerPort: 8081
      name: http
      protocol: TCP
  resources: {}
  restartPolicy: Always
```


- » List a deployment specified in YAML format and redirect it to a file

```
kubectl get deployment <label> -o yaml > file
```

- » This file can be edited and used to create a new deployment


NAMESPACES

Kubernetes allows the ‘physical cluster’ to be split up in multiple virtual clusters. This is done by creating so called `namespaces`. These `namespaces` are different from the Linux kernel `namespaces`. After the K8S cluster has been initialized by `kubeadm` you will find 3 pre-created `namespaces`

- » The `kube-system` namespace containing all Kubernetes infra objects
- » The `kube-public` namespace is specific for `kubeadm`
- » The `default` namespace

NAMESPACES

The standard namespaces available after bootstrapping with `kubeadm`. The `kube-system` is populated with the `k8s-infra` objects



EXAMPLE OF LIST OF PODS RUNNING IN KUBE-SYSTEM NAMESPACE

```
kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-node-qvvvc	2/2	Running	0	124m
coredns-86c58d9df4-zdnhr	1/1	Running	0	126m
etcd-kub14n01	1/1	Running	0	125m
kube-apiserver-kub14n01	1/1	Running	0	125m
kube-controller-manager-kub14n01	1/1	Running	0	125m
kube-proxy-hnsnk	1/1	Running	0	125m
kube-scheduler-kub14n01	1/1	Running	0	125m

- » List namespaces with `kubectl get ns`
- » Create a namespace with `kubectl create ns <namespace>`
- » Delete a namespace with `kubectl delete ns <namespace>`
- » List all pods in a namespace using
`kubectl get pods -n <namespace>`
- » List all pods in all namespaces using
`kubectl get pods --all-namespaces`
- » Create an object in a namespace use
`kubectl create -f mypod.yml -n <namespace>`

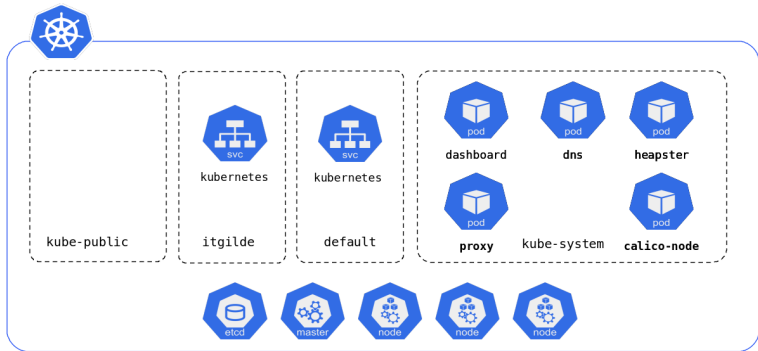
CREATE A NAMESPACE (1)

```
kubectl create ns itgilde
```

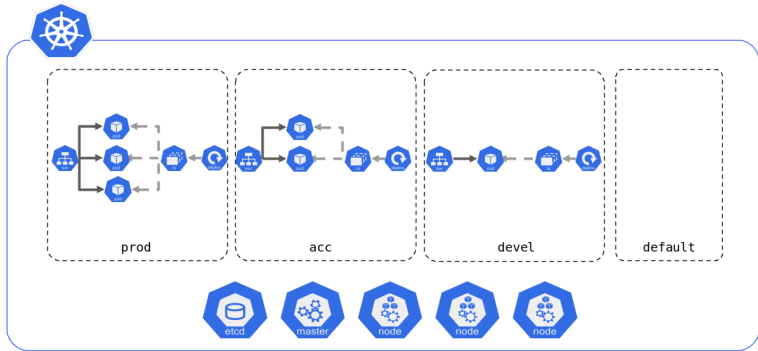
```
kubectl get ns
```

NAME	STATUS	AGE
default	Active	136m
kube-public	Active	136m
kube-system	Active	136m
itgilde	Active	2s

CREATE A NAMESPACE (2)



CREATE A NAMESPACE (3)



- » K8S provides means for Service Discovery using CoreDNS
- » When a service is created to expose a POD a DNS entry is created
- » One can find the POD using:
`<servicename>.<namespace>.svc.cluster.local`
- » If only `<servicename>` is used, it will try to lookup the service local to the namespace
- » Use the FQDN to connect to services across namespaces

OBJECTS THAT ARE NOT LOCAL TO A NAMESPACE

- » namespaces
- » nodes
- » persistentvolumes
- » clusterinformations
- » podsecuritypolicies
- » storageclasses
- » volumeattachments
- » ...

- » persistentvolumeclaims
- » pods
- » replicationcontrollers
- » services
- » daemonsets
- » deployments
- » replicaset
- » ingresses
- » ...

- » Namespaces provide logical partitioning of the Kubernetes cluster
- » There's no TRUE `isolation`
- » Use `namespaces` to separate workloads
- » Use separate clusters to provide `isolation`
- » Alternatives enforced CRI like: `gVisor` Or `Kata Containers`

- » Namespaces are suitable for soft-multitenancy
- » Use it for trusted-workloads within one cluster
- » If you need hard-multitenancy use
 - Separate workloads on separate nodes
 - Enforced CRI
 - Separate workloads on separate clusters

KUBERNETES PODs

KUBERNETES PODS

CONCEPT OF PODS

PODs

- » Hold the containers in K8S
- » Can hold 1 or `multiple` containers
- » Are the unit of scheduling on K8S
- » Get an IP attached to them
- » Get volumes attached to them
- » Are seldom created 'bare'

Containers in a POD

- » Share one IP address
- » Share the attached volumes
- » Can connect to each other using the localhost (127.0.0.1) network

POD manifest in YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: testpod
  labels:
    run: bb
spec:
  containers:
  - args:
    - sleep
    - "3600"
    image: yauritux/busybox-curl
    imagePullPolicy: Always
    name: bb
```

Kubernetes supports 3 type of probes. These probes are configured in the POD spec.

- » `startupProbe`
- » `livenessProbe`
- » `readinessProbe`

- » Monitors the startup of the container
- » When failed: the container is killed
- » When once succesful: disarms startupProbe and arms readinessProbe / livenessProbe

startupProbe

```
apiVersion: v1
kind: Pod
metadata:
  run: bb
spec:
  containers:
    image: pamvdam/astro:sf1
    name: astro
    startupProbe:
      httpGet:
        path: /health
        port: 8080
      failureThreshold: 30
      periodSeconds: 10
```

- » Detects unresponsiveness of the container
- » When failed: the container is killed

livenessProbe

```
apiVersion: v1
kind: Pod
metadata:
  run: bb
spec:
  containers:
    image: pamvdam/astro:sf1
    name: astro
    livenessProbe:
      exec:
        command:
          - cat
          - /tmp/healthy
      initialDelaySeconds: 5
      periodSeconds: 5
```

- » Detects if a container is ready to receive network traffic
- » When failed: network traffic will not be routed to this container

readinessProbe


```
apiVersion: v1
kind: Pod
metadata:
  run: astro
spec:
  containers:
    image: pamvdam/astro:sf1
    name: astro
    readinessProbe:
      exec:
        command:
          - cat
          - /tmp/iamready
      initialDelaySeconds: 5
      periodSeconds: 5
```

There are 4 common Multi-Container/POD Patterns

- » Sidecar Container
- » InitContainer
- » Ambassador Container
- » Adapter Container

MULTI CONTAINER PATTERN 1 - SIDECAR CONTAINER

The `sidecar` pattern allows to extend or augment the functionality of a pre-existing container without changing it.



- » Allow for single-purpose reusable containers
- » Extending the functionality by using Sidecar containers
- » Usecase: initializing an environment for the App containers
- » Usecase: keeping App container config updated (nginx)

Example Sidecar Container

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-sidecar
spec:
  volumes:
    - name: shared-logs
      emptyDir: {}

  containers:
    - name: app-container
      image: alpine
      command: ["/bin/sh"]
      args: ["-c", "while true; do date >> /var/log/app.txt; sleep 5;done"]

    - name: sidecar-container
      image: alpine
      command: ["/bin/sh"]
      args: ["-c", "while true; do date >> /var/log/app.txt; sleep 5;done"]

  volumeMounts:
    - name: shared-logs
      mountPath: /var/log
```


Example Sidecar Container

```
- name: sidecar-container
  image: nginx:1.7.9
  ports:
    - containerPort: 80

  volumeMounts:
    - name: shared-logs
      mountPath: /usr/share/nginx/html
```

MULTI CONTAINER PATTERN 2 - INIT CONTAINER

The `InitContainer` pattern foresees in a means to initialize an environment before the actual application container is run. One could for example clone the most recent website from GIT using an `InitContainer` and once done have it served by an `nginx` container



InitContainers


- » `initContainers` are special 'containers'
- » They have their own 'container' description in the POD manifest
- » They are executed first
- » While `initContainers` are executed, the other containers are not started
- » `initContainers` run to completion
- » If an `initContainer` fails the POD is restarted
- » `initContainers` are started in the order of appearance

InitContainer

```
spec:
  containers:
    - name: web-server
      image: nginx
  initContainers:
    - name: init-clone-repo
      image: alpine/git
      args:
        - clone
        - --single-branch
        - --
        - https://thegitcave.org/k8s4all/website.repo
        - /usr/share/nginx/html
```

MULTI CONTAINER PATTERN 3 - AMBASSADOR CONTAINER


The Ambassador Container pattern is a specialized `Sidecar` pattern that provides a unified interface for accessing services outside of the pod



- » Does not enhance the app container
- » Provides the function of a smartcache
- » Usecase: SSL termination
- » Usecase: memcached and twemproxy

MULTI CONTAINER PATTERN 4 - ADAPTER CONTAINER

The Adapter Container pattern provides a way to have uniform access to a heterogeneous system.



- » Usecase: Expose metrics in a standard way
- » Usecase: Expose log records in a standard way
- » Actually a specialized case of the `Sidecar` pattern
- » Or a reverse `Ambassador` pattern

kubectl and PODs

To look inside a POD, docker exec equivalent

```
kubectl exec -it <podname> [-c <containername>] -- sh
```

To execute a command inside a pod

```
kubectl exec <podname> [-c <containername>] -- cat /etc/hosts
```

To get the logs from a container in a pod

```
kubectl logs <podname> [-c <containername>]
```

To follow the logs from a container in a pod

```
kubectl logs -f <podname> [-c <containername>]
```

KUBERNETES DEPLOYMENTS

KUBERNETES DEPLOYMENTS

CONCEPT OF DEPLOYMENTS

A Deployment

- » Provides a way to deploy managed `ReplicaSets`
- » The generated `ReplicaSet` will deploy a set of identical `PODs`
- » Gives a more declarative interface to `RS` and `POD` updates
- » The `DeploymentController` manages the `Deployment`

With Deployments

- » One can deploy a `ReplicaSet`
- » One can update PODs
- » One can rollback to older versions of the `Deployment`
- » One can pause and resume the `Deployment`
- » One can execute various `Deployment` and upgrade patterns

To create a Deployment adhoc using `kubectl` execute:

```
kubectl create deployment nginx --image nginx:1.7.1
```

To create a Deployment using a YAML manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: httpd
  replicas: 4
  template:
    metadata:
      labels:
        app: httpd
    spec:
      containers:
        - name: apache
          image: httpd:2.4.39-alpine
          ports:
            - containerPort: 80
```

The YAML manifest of a Deployment has the following important sections and attributes:

- » `spec.replicas` defines the nr of replicas
- » `spec.selector.matchLabels` PODs with this label will be managed by the Deployment
- » `spec.template` This defines what kind of PODs need to be (re-)created

Tasks on Deployments

```
# To view deployments in the K8S cluster
kubectl get deployments
kubectl get deploy

# To describe the state of a deployment
kubectl describe deployment httpd

# To delete a deployment
kubectl delete deployment httpd

# To update a deployment
kubectl set image deployment/nginx-deployment nginx=nginx:1.91 --record

# To display the history of updates
kubectl rollout history deployment.v1.apps/nginx-deployment

# To rollback an update
kubectl rollout undo deployment.v1.apps/nginx-deployment

# List PODs created by this deployment
kubectl get pods -l httpd
```

KUBERNETES DAEMONSETS

KUBERNETES DAEMONSETS

CONCEPT OF DAEMONSETS

Sometimes a POD only needs to run one single instance per node. The DaemonSet ensures this. A `DaemonSet`

- » Manages a set of PODs
- » Ensures that each `node` gets exactly 1 POD
- » When a `node` is added to the cluster this `node` will automatically get it's own instance of the POD
- » When a `node` is removed from the cluster, no other node will receive an extra POD.

The following cases are suitable for deploying PODs using `DaemonSets`

- » Security, vulnerability or virus scanners
- » Logging agents
- » Performance collector agents
- » Ingress controllers

DaemonSets are created with YAML manifests:

```
kind: DaemonSet
metadata:
  name: fluentd-elasticsearch
  labels:
    k8s-app: fluentd-logging
spec:
  selector:
    matchLabels:
      name: fluentd-elasticsearch
  template:
    metadata:
      labels:
        name: fluentd-elasticsearch
    spec:
      containers:
        - name: fluentd-elasticsearch
          image: k8s.gcr.io/fluentd-elasticsearch:1.20
```

The YAML manifest of a `DaemonSet` has the following important sections and attributes:

- » `spec.selector.matchLabels` PODs with this label will be managed by the `DaemonSet`
- » `spec.template` This defines what kind of PODs need to be (re-)created

Tasks on DaemonSets

To view daemonsets in the K8S cluster

```
kubectl get daemonsets
```

```
kubectl get ds
```

To describe the state of a daemonset

```
kubectl describe ds frontend
```

To delete a daemonset

```
kubectl delete ds frontend
```

You may now work on the LABs in chapter 5

- » 5.1 Creating DaemonSets
- » 5.2 Communicating with PODs managed by DaemonSets
- » 5.3 Upgrading PODs in DaemonSets

KUBERNETES JOBS

KUBERNETES JOBS

CONCEPT OF JOBS

PODs can be restarted automatically upon exiting using a `ReplicaSet` or `Deployment`. This ideal for PODs that have to process an (virtually) infinite amount of work. However some PODs have work that is finite and only need to be restarted upon failure and not upon completion. For these type of PODs K8S provides the `Job`. These `Jobs`

- » Manage a set of PODs to carry out a finite amount of workload
- » Will restart them upon failure
- » Will not restart them upon completion

Use cases for Jobs

- » Batch processing of a finite amount of work at a time
- » Work that needs to be done to transform or update data
- » Work that setups an environment for other PODs

To create a Job adhoc using `kubectl` execute:

```
kubectl create job busybox --image=busybox
```

To create a Job using a YAML manifest

```
apiVersion: batch/v1
kind: Job
metadata:
  name: example-job
spec:
  template:
    metadata:
      name: example-job
    spec:
      containers:
      - name: pi
        image: perl
        command: ["perl"]
        args: ["-Mbignum=bpi", "-wle", "print bpi(2000)"]
      restartPolicy: Never
```

The YAML manifest of a `Job` has the following important sections and attributes:

- » `spec.template.spec.restartPolicy` Always set to `Never` for a `Job`
- » `spec.template` This defines what kind of `POD` needs to be (re-)created

Tasks on Jobs

To view jobs in the K8S cluster

```
kubectl get jobs
```

To describe the state of a job

```
kubectl describe job example-job
```

To delete a job

```
kubectl delete job httpd
```

KUBERNETES CRONJOBS

KUBERNETES CRONJOBS

CONCEPT OF CRONJOBS

If a Job needs to be executed at a scheduled time, K8S provides so called `CronJobs` for this purpose. `CronJobs`:

- » Execute a set of PODs using a predefined schedule
- » Use a UNIX/Linux `crontab` like notation
- » Jobs that are completed are not restarted
- » Jobs that fail get restarted

Use cases for CronJobs

- » Batch processing of a finite amount of work periodically
- » End of day processing
- » Periodic scanning

To create a CronJob adhoc using `kubectl` execute:

```
kubectl create cronjob busybox --image=busybox --schedule="4 10 * * *"
```

To create a CronJob using a YAML manifest

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: my-crontab
spec:
  schedule: "*/5 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: pi
              image: perl
              command: ["perl"]
              args: ["-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: OnFailure
```

The YAML manifest of a `CronJob` has the following important sections and attributes:

- » `spec.schedule` Specifies when the Job should be scheduled
- » `spec.jobTemplate` Describes the job to be executed

Tasks on Jobs

To view cronjobs in the K8S cluster

```
kubectl get cronjobs
```

```
kubectl get cj
```

To describe the state of a cronjob

```
kubectl describe cronjob my-cronjob
```

To delete a cronjob

```
kubectl delete cronjob my-cronjob
```

KUBERNETES PERSISTENT STORAGE

- » Like Docker containers, pods are designed to be volatile
- » This means they cannot keep state themselves
- » If an application needs to keep state, you'll need (persistent) storage

- » Local storage
- » iSCSI
- » NFS
- » Cloud storage
- » emptyDir
- » hostMount
- » configMaps
- » secrets
- » etc...

- » Persistence: only for the lifetime of the POD
- » Can be shared with other containers in the POD

```
kind: Pod
apiVersion: v1
metadata:
  name: simple-volume-pod
spec:
  volumes:
    - name: simple-vol
      emptyDir: {}
  containers:
    - name: my-container
      volumeMounts:
        - name: simple-vol
          mountPath: /var/simple
      image: alpine
      command: ["/bin/sh"]
      args: ["-c", "while true; do date >> /var/simple/file.txt; sleep 5; done"]
```


- » Persistent Volume (PV) resources are used to manage durable storage in a cluster
- » Unlike volumes the lifecycle is managed by Kubernetes
- » A PV can already exist or dynamically provisioned via plugins
- » The PV must be bind to the cluster using a Persistent Volume Claim (PVC)
- » The PVC dictates the kind and size of the storage

- » Create one or more PVs (they map storage)
- » Create a Persistent Volume Claim
- » A POD is created claiming storage using a PVC
- » The scheduler selects which PV is suitable for the PVC
- » Storage is bound to the POD

PERSISTENT VOLUME MANIFEST (1)

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-nfs-002
spec:
  capacity:
    storage: 20Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
```

```
mountOptions:  
  - hard  
  - nfsvers=4.1  
nfs:  
  path: /k8s/vol002  
  server: 10.8.62.222
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-nfs-001
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: "slow"
  resources:
    requests:
      storage: 1Mi
```

```
spec:
  containers:
  - name: pod-pvc
    ports:
    ...
    volumeMounts:
      # name must match the volume name below
      - name: nfs-html
        mounthPath: "/usr/share/nginx/html"
  volumes:
  - name: nfs-html
    persistentVolumeClaim:
      claimName: pvc-nfs-001
```

KUBERNETES CONFIGMAPS

- » Like Docker configs, ConfigMaps are used to store configuration data
- » They are used to separate the configuration from the container image
- » ConfigMaps are not encrypted

ConfigMaps can materialize the data contained as:

- » Environment Variables in the container/POD
- » Files in a filesystem in the container/POD

Adhoc from a file

```
kubectl create configmap mysqlcfg1 --from-file=/etc/mysql.conf
```

Adhoc from literal

```
kubectl create configmap myconfig --from-literal=color=red --from-literal=mascot=astro
```

```
apiVersion: v1
data:
  color: red
  mascot: astro
kind: ConfigMap
metadata:
  name: myconfig
```

```
apiVersion: v1
kind: Pod
metadata:
  name: color-container
spec:
  containers:
    - name: color-container
      image: pamvdam/nginx:v1.1
      env:
        - name: COLOR
          valueFrom:
            configMapKeyRef:
              name: myconfig
              key: color
      restartPolicy: Never
```

```
apiVersion: v1
kind: Pod
metadata:
  name: color-container
spec:
  containers:
    - name: color-container
      image: pamvdam/nginxc:v1.1
      volumemounts:
        - name: myconfigvol
          mountPath: /myconfig
  volumes:
    - name: myconfigvol
      configMap:
        name: myconfig
```

Delete a configmap

```
kubectl delete configmap myconfig
```

Describe a configmap

```
kubectl describe configmap myconfig
```

Update PODs in a deployment where ConfigMap has been updated

```
kubectl rollout restart deploy
```

You may now start with LAB 7 - ConfigMaps

KUBERNETES SECRETS

- » Like Docker secrets, Secrets are used to store sensitive configuration data
- » They are used to separate this type of configuration from the container image
- » Use secrets to store privkeys, passwords, certs etc
- » Secrets are encoded not encrypted
- » Use encryption at rest to get them encrypted

Secrets can materialize the data contained as:

- » Environment Variables in the container/POD
- » Files in a filesystem in the container/POD

Adhoc from a file

```
kubectl create secret generic mysecret1 --from-file=/etc/mysql.passwd
```

Adhoc from literal

```
kubectl create secret generic mysecret2 --from-literal=color=red --from-literal=mascot=astro
```

```
apiVersion: v1
kind: Secret
metadata:
  name: test-secret
data:
  username: bXktYXBw
  password: Mzk1MjgkdmRnNOpi
```

```
apiVersion: v1
kind: Pod
metadata:
  name: env-single-secret
spec:
  containers:
  - name: envvars-test-container
    image: nginx
    env:
    - name: SECRET_PASSWORD
      valueFrom:
        secretKeyRef:
          name: test-secret
          key: password
```

USING SECRETS IN PODS - VOLUMEMOUNT

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-test-pod
spec:
  containers:
    - name: test-container
      image: nginx
      volumeMounts:
        - name: secret-volume
          mountPath: /etc/secret-volume
  volumes:
    - name: secret-volume
      secret:
        secretName: test-secret
```

How to pull an image in a POD from a `authenticated registry`

- » Create an `imagePullSecrets`
- » Have the POD consume the `imagePullSecrets`

```
kubectl create secret docker-registry my-pullsecret \  
  --docker-server=<your-registry-server> \  
  --docker-username=<your-name> \  
  --docker-password=<your-pword> \  
  --docker-email=<your-email
```



```
apiVersion: v1
kind: Pod
metadata:
  name: private-reg
spec:
  containers:
    - name: private-reg-container
      image: harbor-prod.itgildelab.net/pyco2:unreleasedv6
  imagePullSecrets:
    - name: my-pullsecret
```

Delete a secret

```
kubectl delete secret mysecret
```

Describe a secret

```
kubectl describe secret mysecret
```

Update PODs in a deployment where Secret has been updated

```
kubectl rollout restart deploy
```